

# Linear Regression and Resampling Methods for Uncertainty

## Big Data y Machine Learning para Economía Aplicada

Ignacio Sarmiento-Barbieri

Universidad de los Andes

# Agenda

- 1 Review
- 2 Uncertainty: Motivation
  - The Bootstrap
    - Example: Elasticity of Demand for Gasoline
- 3 Train and Test Sets. In-Sample and Out-of-Sample Prediction.
  - AIC: Akaike Information Criterion
  - SIC/BIC: Schwarz/Bayesian Information Criterion
  - Cross-Validation
- 4 Review

# Agenda

## 1 Review

## 2 Uncertainty: Motivation

- The Bootstrap
  - Example: Elasticity of Demand for Gasoline

## 3 Train and Test Sets. In-Sample and Out-of-Sample Prediction.

- AIC: Akaike Information Criterion
- SIC/BIC: Schwarz/Bayesian Information Criterion
- Cross-Validation

## 4 Review

# Predicting Well

$$y = f(X) + u \quad (1)$$

- ▶ Interest on predicting  $y$
- ▶ Under quadratic loss  $\Rightarrow E[y|X = x]$

# Linear Regression

$$y = f(X) + u \quad (2)$$

$$= \beta_0 + \beta_1 X_1 + \cdots + \beta_k X_k + u \quad (3)$$

$$= X\beta + u \quad (4)$$

- If  $f(X) = X\beta$ , obtaining  $f(\cdot)$  boils down to obtaining  $\beta$

# Linear Regression

- ▶ OLS says we should choose the estimators  $\hat{\beta}$  such that we minimize the Sum of Square Residual (SSR)

$$\mathcal{L} = \sum_{i=1}^n (y_i - \hat{y}_i)^2 \quad (5)$$

$$= \sum_{i=1}^n \left( y_i - \hat{\beta}_0 - \sum_{j=1}^k \hat{\beta}_j x_{ji} \right)^2 \quad (6)$$

$$= (y - X\hat{\beta})'(y - X\hat{\beta}) \quad (7)$$

- ▶ Compute  $\beta$ 
  - ▶ QR (Gram-Schmidt process, similar to FWL)
  - ▶ Gradient Descent

# Agenda

## 1 Review

## 2 Uncertainty: Motivation

- The Bootstrap

- Example: Elasticity of Demand for Gasoline

## 3 Train and Test Sets. In-Sample and Out-of-Sample Prediction.

- AIC: Akaike Information Criterion
- SIC/BIC: Schwarz/Bayesian Information Criterion
- Cross-Validation

## 4 Review

# Uncertainty in Linear Regression

- ▶ To get a measure of the uncertainty, precision or variability of our estimates we need a measure
- ▶ We can estimate the Variance of our estimators
- ▶ Linear regression

$$\text{Var}(\hat{\beta}) = \text{Var}((X'X)^{-1}X'y) \quad (8)$$



# Uncertainty and Resampling

- ▶ Sometimes the analytical expression of the variance can be quite complicated.
- ▶ In these cases we can use the bootstrap
- ▶ The bootstrap provides a way to perform statistical inference by resampling from the sample.

# Agenda

## 1 Review

## 2 Uncertainty: Motivation

- The Bootstrap

- Example: Elasticity of Demand for Gasoline

## 3 Train and Test Sets. In-Sample and Out-of-Sample Prediction.

- AIC: Akaike Information Criterion
- SIC/BIC: Schwarz/Bayesian Information Criterion
- Cross-Validation

## 4 Review

# The Bootstrap

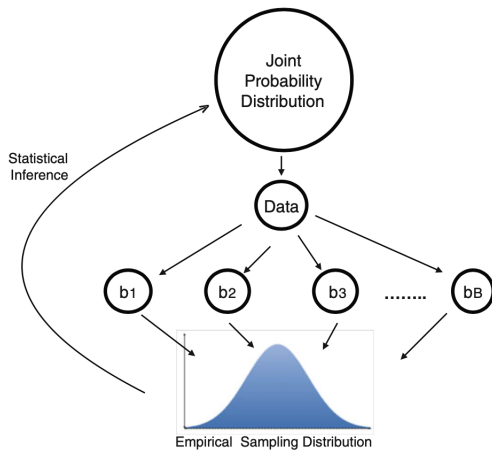
## Introduction

- ▶ The bootstrap is a widely applicable and extremely powerful statistical tool that can be used to quantify the uncertainty associated with a given estimator or statistical learning method.
- ▶ In German the expression *an den eigenen Haaren aus dem Sumpf zu ziehen* nicely captures the idea of the bootstrap – “to pull yourself out of the swamp by your own hair.”



# The Bootstrap

## Introduction



# The Bootstrap

## Introduction

- ▶ There are two key properties of bootstrapping that make this seemingly crazy idea actually work.
  - 1 Each bootstrap sample must be of the same size ( $N$ ) as the original sample
  - 2 Each bootstrap sample must be taken with replacement from the original sample.

# The Bootstrap

- ▶ In general terms:
  - ▶  $Y_i$   $i = 1, \dots, n$
  - ▶  $\theta$  is the magnitude of interest
- ▶ To calculate it's variance
  - 1 Sample of size  $n$  with replacement (*bootstrap sample*)
  - 2 Compute  $\hat{\theta}_j$   $j = 1, \dots, B$
  - 3 Repeat  $B$  times
  - 4 Calculate

$$\hat{V}(\hat{\theta})_B = \frac{1}{B} \sum_{j=1}^B (\hat{\theta}_j - \bar{\hat{\theta}})^2 \quad (9)$$

# Example: Elasticity of Demand for Gasoline



photo from <https://www.dailydot.com/parsec/batman-1966-labels-tumblr-twitter-vine/>

# Agenda

## ① Review

## ② Uncertainty: Motivation

- The Bootstrap
  - Example: Elasticity of Demand for Gasoline

## ③ Train and Test Sets. In-Sample and Out-of-Sample Prediction.

- AIC: Akaike Information Criterion
- SIC/BIC: Schwarz/Bayesian Information Criterion
- Cross-Validation

## ④ Review



# Train and Test Sets. In-Sample and Out-of-Sample Prediction.

- ▶ El objetivo es predecir  $y$  dadas otras variables  $X$ . Ej: salario dadas las características del individuo
- ▶ Asumimos que el link entre  $y$  and  $X$  esta dado por el modelo:

$$y = f(X) + u \quad (10)$$

- ▶ donde  $f(X)$  por ejemplo es  $\beta_0 + \beta_1 X_1 + \dots + \beta_k X_k$
- ▶  $u$  una variable aleatoria no observable  $E(u) = 0$  and  $V(u) = \sigma^2$

# Train and Test Sets. In-Sample and Out-of-Sample Prediction.

- ▶ Dos conceptos importantes
  - ▶ *Training error*: es el error de predicción en la muestra que fue utilizada para ajustar el modelo

$$Err_{\mathcal{T}_{rain}} = MSE[(y, \hat{y}) | \mathcal{T}_{rain}] \quad (11)$$

- ▶ *Test Error*: es el error de predicción fuera de muestra

$$Err_{\mathcal{T}_{est}} = MSE[(y, \hat{y}) | \mathcal{T}_{est}] \quad (12)$$

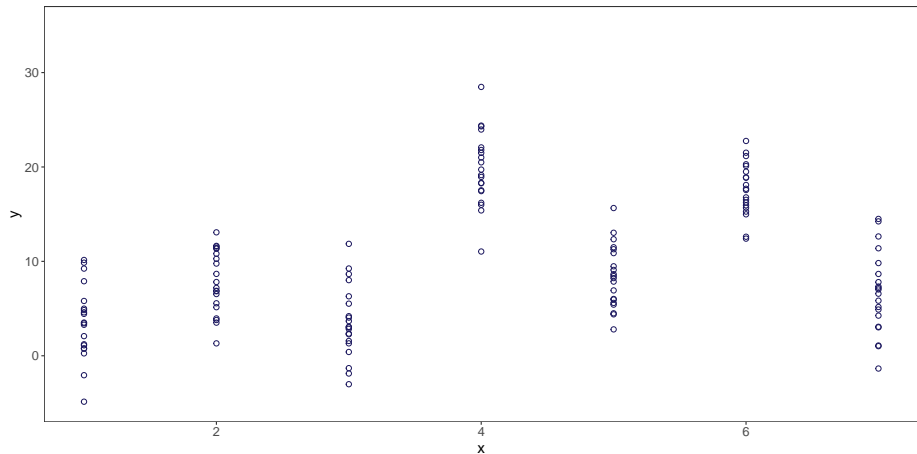
# Train and Test Sets. In-Sample and Out-of-Sample Prediction.

- Como seleccionamos la especificación que minimize el error de predicción?

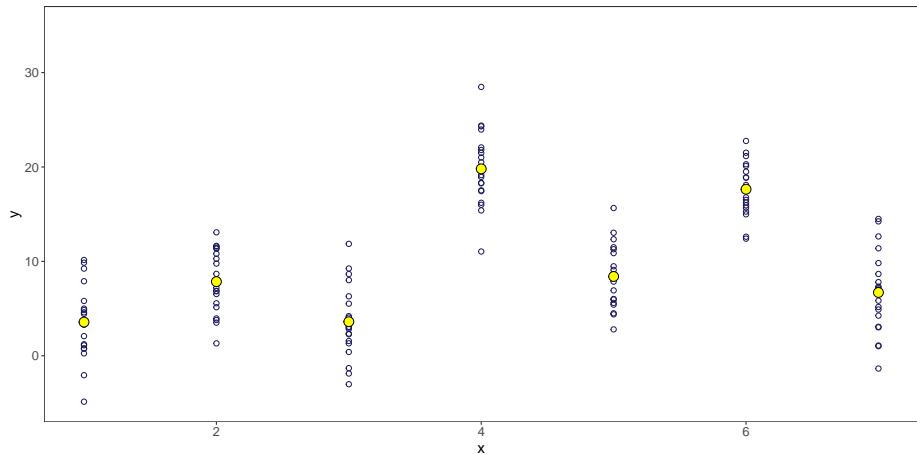
# Train and Test Sets. In-Sample and Out-of-Sample Prediction.

- ▶ Como seleccionamos la especificación que minimize el error de predicción?
- ▶ Problema: solo contamos con una muestra

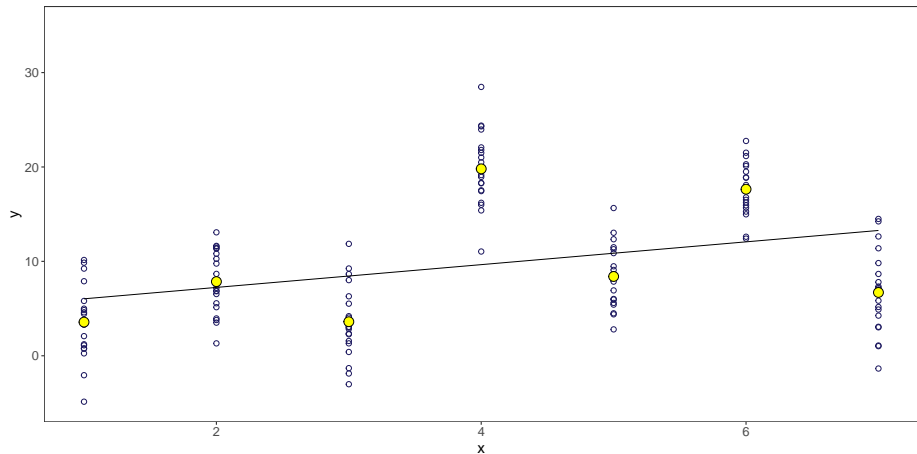
# In-Sample Prediction and Overfit



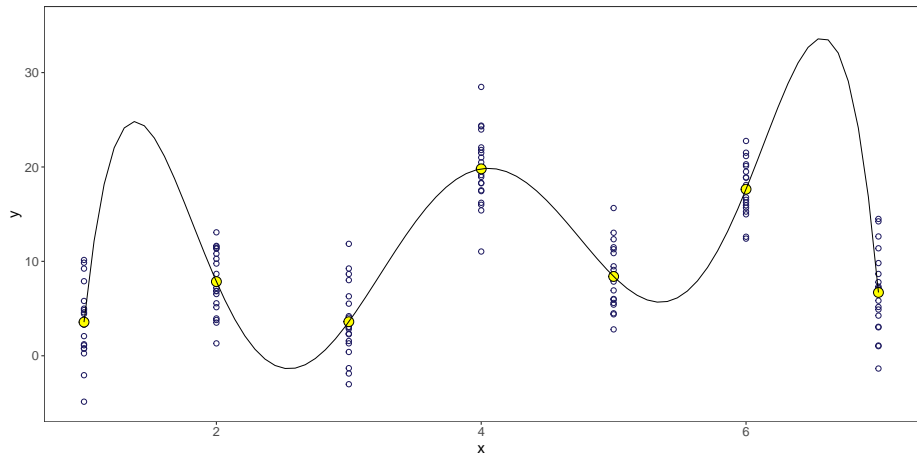
# In-Sample Prediction and Overfit



# In-Sample Prediction and Overfit

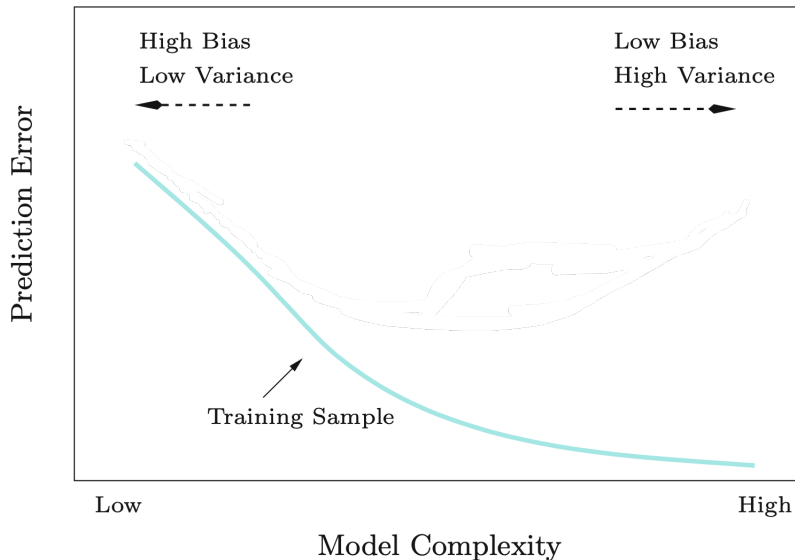


# In-Sample Prediction and Overfit





# In-Sample Prediction and Overfit



# In-Sample Prediction and Overfit

- Notemos que el MSE no es otra cosa que la suma de los residuales al cuadrado

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{f}(X))^2 \quad (13)$$

$$= \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y})^2 \quad (14)$$

$$= \frac{1}{n} \sum_{i=1}^n (e)^2 \quad (15)$$

$$= SSR \quad (16)$$

- Esta medida nos da una idea de *lack of fit* que tan mal ajusta el modelo a los datos

# In-Sample Prediction and Overfit

- ▶ Un problema del SSR es que nos da una medida absoluta de ajuste de los datos, y por lo tanto no está claro que constituye un buen SRR.
- ▶ Una alternativa muy usada en economía es el  $R^2$
- ▶ Este es una proporción (la proporción de varianza explicada),
  - ▶ toma valores entre 0 y 1,
  - ▶ es independiente de la escala (o unidades) de  $y$

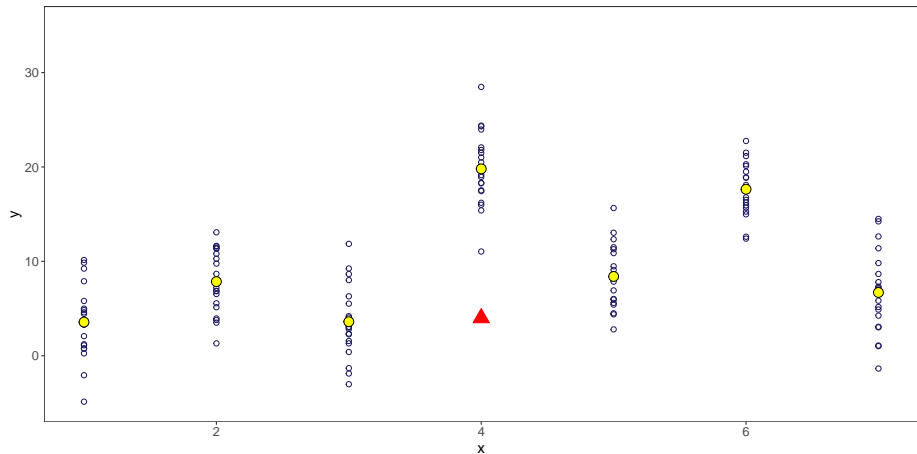
$$R^2 = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y})^2}{\sum_{i=1}^n (y_i - \bar{y})^2} \quad (17)$$

$$= 1 - \frac{SRR}{TSS} \quad (18)$$

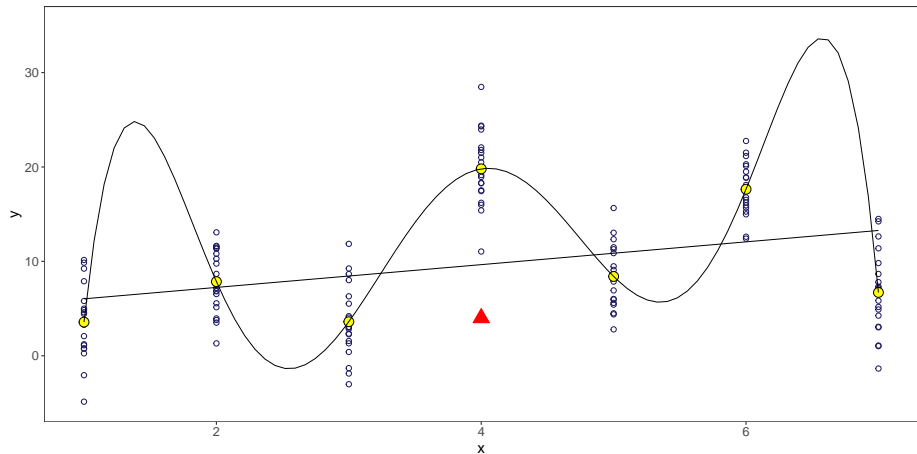
# Out-of-Sample Prediction and Overfit

- ▶ ML nos interesa la predicción fuera de muestra

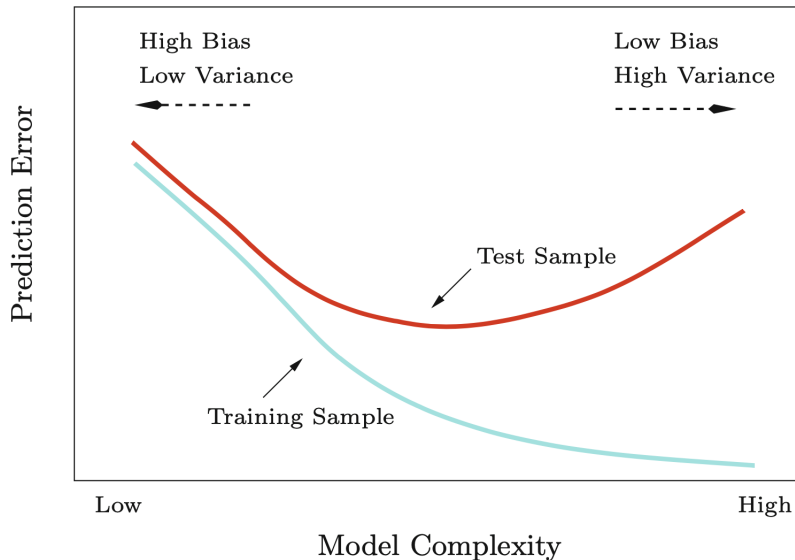
# Out-of-Sample Prediction and Overfit



# Out-of-Sample Prediction and Overfit



# Out-of-Sample Prediction and Overfit



# Out-of-Sample Prediction and Overfit

- ▶ ML nos interesa la predicción fuera de muestra
- ▶ Overfit: modelos complejos predicen muy bien dentro de muestra, pero tienden a hacer un trabajo fuera de muestra
- ▶ Hay que elegir el nivel adecuado de complejidad
- ▶ Como medimos el error de predicción fuera de muestra?
- ▶  $R^2$  no funciona: se concentra en la muestra y es no decreciente en complejidad



# Test Error

- ▶ Para seleccionar el mejor modelo con respecto al Test Error (error de prueba), es necesario estimarlo.
- ▶ Hay dos enfoques comunes:
  - ▶ Podemos estimar indirectamente el error de la prueba haciendo un ajuste al error de entrenamiento para tener en cuenta el sesgo debido al sobreajuste  $\Rightarrow$  Penalización ex post: AIC, BIC, R2 ajustado

# Agenda

## 1 Review

## 2 Uncertainty: Motivation

- The Bootstrap
  - Example: Elasticity of Demand for Gasoline

## 3 Train and Test Sets. In-Sample and Out-of-Sample Prediction.

- AIC: Akaike Information Criterion
- SIC/BIC: Schwarz/Bayesian Information Criterion
- Cross-Validation

## 4 Review

# Test Error

## AIC

- ▶ Akaike (1969) fue el primero en ofrecer un enfoque unificado al problema de la selección de modelos.
- ▶ Elegir el modelo  $j$  tal que se minimice:

$$AIC(j) = \log \left( \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y})^2 \right) - p_j \quad (19)$$

# Agenda

## 1 Review

## 2 Uncertainty: Motivation

- The Bootstrap
  - Example: Elasticity of Demand for Gasoline

## 3 Train and Test Sets. In-Sample and Out-of-Sample Prediction.

- AIC: Akaike Information Criterion
- SIC/BIC: Schwarz/Bayesian Information Criterion
- Cross-Validation

## 4 Review

# Test Error

## SIC/BIC

- ▶ Schwarz (1978) mostró que el AIC es inconsistente, (cuando  $n \rightarrow \infty$ , tiende a elegir un modelo demasiado grande con probabilidad positiva)
- ▶ Schwarz (1978) propuso:

$$SIC(j) = \log \left( \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y})^2 \right) - \frac{1}{2} p_j \log(n) \quad (20)$$

# Test Error

## AIC vs BIC

$$AIC(j) = \log \left( \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y})^2 \right) - p_j \quad (21)$$

$$SIC(j) = \log \left( \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y})^2 \right) - p_j \frac{1}{2} \log(n) \quad (22)$$

- ▶ SIC tiende a elegir modelos más pequeños.
- ▶ En efecto, al dejar que la penalización tienda al infinito lentamente con  $n$ , eliminamos la tendencia de AIC a elegir un modelo demasiado grande.

# Test Error

- ▶ Para seleccionar el mejor modelo con respecto al Test Error (error de prueba), es necesario estimarlo.
- ▶ Hay dos enfoques comunes:
  - ▶ Podemos estimar indirectamente el error de la prueba haciendo un ajuste al error de entrenamiento para tener en cuenta el sesgo debido al sobreajuste  $\Rightarrow$  Penalización ex post: AIC, BIC,  $R^2$  ajustado
  - ▶ Levantarnos de nuestros bootstraps (resampling methods) y estimar directamente el Test Error (error de prueba)

# Agenda

## 1 Review

## 2 Uncertainty: Motivation

- The Bootstrap
  - Example: Elasticity of Demand for Gasoline

## 3 Train and Test Sets. In-Sample and Out-of-Sample Prediction.

- AIC: Akaike Information Criterion
- SIC/BIC: Schwarz/Bayesian Information Criterion
- Cross-Validation

## 4 Review



# Test Error

## Cross-Validation



photo from <https://www.dailydot.com/parsec/batman-1966-labels-tumblr-twitter-vine/>

# Agenda

## ① Review

## ② Uncertainty: Motivation

- The Bootstrap
  - Example: Elasticity of Demand for Gasoline

## ③ Train and Test Sets. In-Sample and Out-of-Sample Prediction.

- AIC: Akaike Information Criterion
- SIC/BIC: Schwarz/Bayesian Information Criterion
- Cross-Validation

## ④ Review

# Review

Hoy

- ▶ Dilema Sesgo/Varianza
- ▶ Sobreajuste y Selección de modelos
  - ▶ AIC y BIC
  - ▶ Enfoque de Validación
  - ▶ LOOCV
  - ▶ K-fold Cross-Validation (Validación Cruzada)