

Generative Models for Classification and Missclassification

Big Data y Machine Learning para Economía Aplicada

Ignacio Sarmiento-Barbieri

Universidad de los Andes

Agenda

- 1 Recap
- 2 Generative Models for Classification
 - Discriminant Analysis
 - Naive Bayes
- 3 Misclassification Rates
 - ROC curve
- 4 Multiple Classes
 - KNN
 - Multinomial Logit

Agenda

- 1 Recap
- 2 Generative Models for Classification
 - Discriminant Analysis
 - Naive Bayes
- 3 Misclassification Rates
 - ROC curve
- 4 Multiple Classes
 - KNN
 - Multinomial Logit

Recap

- ▶ We observe (y_i, X_i) $i = 1, \dots, n$
- ▶ Estimate Probabilities
 - ▶ Logit

$$p_i = \frac{e^{X_i \beta}}{1 + e^{X_i \beta}} \quad (1)$$

- ▶ get β
- ▶ Prediction
 - ▶ Logit, with the $\hat{\beta}$

$$\hat{p}_i = \frac{e^{X_i \hat{\beta}}}{1 + e^{X_i \hat{\beta}}} \quad (2)$$

- ▶ Classification

$$\hat{Y}_i = 1[\hat{p}_i > c] \quad (3)$$


Agenda

- ① Recap
- ② Generative Models for Classification
 - Discriminant Analysis
 - Naive Bayes
- ③ Misclassification Rates
 - ROC curve
- ④ Multiple Classes
 - KNN
 - Multinomial Logit

Linear Discriminant Analysis

Reverend Bayes to the rescue: Bayes Theorem

$$Pr(Y = 1|X) \quad (4)$$



Linear Discriminant Analysis, $k=1$

Linear Discriminant Analysis, $k \geq 1$

Extensions

- ▶ If we have $k > 1$ predictors?
- ▶ then $X|Y \sim NM(\mu, \Sigma)$

$$f(X|Y = j) = \frac{1}{(2\pi)^{k/2} |\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(x - \mu_j)' \Sigma_j (x - \mu_j)\right) \quad (5)$$

- ▶ μ_j is the vector of the sample means in each partition $j = 0, 1$
- ▶ Σ_j is the matrix of variance and covariances of each partition $j = 0, 1$

Example: Default





photo from <https://www.dailydot.com/parsec/batman-1966-labels-tumblr-twitter-vine/>

Linear Discriminant Analysis

- ▶ Why is it called linear?
- ▶ One predictor with $\sigma_0 = \sigma_1$ (equal variance)

Quadratic Discriminant Analysis

- QDA assumes different variances for the components



Naive Bayes


$$Pr(Y = 1|X) = \frac{f(X|Y = 1)\pi(Y = 1)}{f(X|Y = 1)\pi(Y = 1) + f(X|Y = 0)(1 - \pi(Y = 1))} \quad (6)$$

- ▶ $\pi(Y = 1)$
- ▶ $f(X|Y = 1)$

Naive Bayes

- NB assumes independence

$$f(X|Y = 1) = f(x_1|Y = 1) \times \cdots \times f(x_k|Y = 1) \quad (7)$$



Example: Default




photo from <https://www.dailydot.com/parsec/batman-1966-labels-tumblr-twitter-vine/>

Agenda

- 1 Recap
- 2 Generative Models for Classification
 - Discriminant Analysis
 - Naive Bayes
- 3 Misclassification Rates
 - ROC curve
- 4 Multiple Classes
 - KNN
 - Multinomial Logit

Misclassification Rates



Misclassification Rates

		\hat{y}_i	
		0	1
y_i	0	TN	FP
	1	FN	TP

- We have several types of error associated with this that we can use as a measure of performance

ROC


		\hat{y}_i	
		0	1
y_i	0	TN	FP
	1	FN	TP

- ▶ A classification rule, or cutoff, is the probability p at which you predict
 - ▶ $\hat{y}_i = 0$ if $p_i < c$
 - ▶ $\hat{y}_i = 1$ if $p_i > c$
- ▶ Bayes classifier $c = 0.5$
- ▶ Changing c changes predictions, changes FP and FN
- ▶ There is a trade-off: reducing one error increases the other


ROC

- ▶ ROC curve: Receiver operating characteristic curve
- ▶ ROC curve illustrates the trade-off of the classification rule
- ▶ Gives us the ability
 - ▶ Measure the predictive capacity of our model
 - ▶ Compare between models

ROC



ROC



Example: Default




photo from <https://www.dailydot.com/parsec/batman-1966-labels-tumblr-twitter-vine/>

Agenda

- 1 Recap
- 2 Generative Models for Classification
 - Discriminant Analysis
 - Naive Bayes
- 3 Misclassification Rates
 - ROC curve
- 4 Multiple Classes
 - KNN
 - Multinomial Logit

K-Nearest Neighbors

- ▶ What happens when we have to predict multiple outcomes?
- ▶ K nearest neighbor (K-NN) algorithm predicts class \hat{y} for x by asking *What is the most common class for observations around x ?*



K-Nearest Neighbors

- ▶ K nearest neighbor (K-NN) algorithm predicts class \hat{y} for x by asking *What is the most common class for observations around x ?*
- ▶ Algorithm: given an input vector x_f where you would like to predict the class label
 - ▶ Find the K nearest neighbors in the dataset of labeled observations, $\{x_i, y_i\}_{i=1}^n$, the most common distance is the Euclidean distance:

$$d(x_i, x_f) = \sqrt{\sum_{j=1}^p (x_{ij} - x_{fj})^2} \quad (8)$$

- ▶ This yields a set of the K nearest observations with labels:

$$[x_{i1}, y_{i1}], \dots, [x_{iK}, y_{iK}] \quad (9)$$

- ▶ The predicted class of x_f is the most common class in this set

$$\hat{y}_f = \text{mode}\{y_{i1}, \dots, y_{iK}\} \quad (10)$$

K-Nearest Neighbors


- There are some major problems with practical implications
 - Knn predictions are unstable as a function of K

$$K = 1 \implies \hat{p}(\text{white}) = 0$$

$$K = 2 \implies \hat{p}(\text{white}) = 1/2$$

$$K = 3 \implies \hat{p}(\text{white}) = 2/3$$

$$K = 4 \implies \hat{p}(\text{white}) = 1/2$$



Source: Taddy (2019)

K-Nearest Neighbors

- ▶ There are some major problems with practical implications
 - ▶ Knn predictions are unstable as a function of K
 - ▶ This instability of prediction makes it hard to choose the optimal K and cross validation doesn't work well for KNN
 - ▶ Since prediction for each new x requires a computationally intensive counting, KNN is too expensive to be useful in most big data settings.
 - ▶ KNN is a good idea, but too crude to be useful in practice

The multinomial logit model: Intuition

- ▶ The MNLM can be thought of as simultaneously fitting binary logits for all comparisons among the alternatives.
- ▶ For example,
 - ▶ We have a categorical variable with the outcomes for Democrat, for Independent, and for Republican.
 - ▶ Assume that there is one independent variable measuring income in 1,000s.
 - ▶ We can examine the effect of income on party by fitting three binary logits,

The multinomial logit model: Intuition

- ▶ The MNLM can be thought of as simultaneously fitting binary logits for all comparisons among the alternatives.

$$\ln \frac{Pr(D|X)}{Pr(I|X)} = \beta_{0,D|I} + \beta_{1,D|I}Income \quad (11)$$

$$\ln \frac{Pr(R|X)}{Pr(I|X)} = \beta_{0,R|I} + \beta_{1,R|I}Income$$

$$\ln \frac{Pr(D|X)}{Pr(R|X)} = \beta_{0,D|R} + \beta_{1,D|R}Income$$

- ▶ where the subscripts to the β 's indicate which comparison is being made.

The multinomial logit model: Intuition

- These logits include redundant info

$$\ln \frac{Pr(D|X)}{Pr(I/X)} - \ln \frac{Pr(R|X)}{Pr(I/X)} = \ln \frac{Pr(D|X)}{Pr(R/X)} \quad (12)$$

- which implies

$$\beta_{0,D|I} - \beta_{0,R|I} = \beta_{0,D|R} \quad (13)$$

$$\beta_{1,D|I} - \beta_{1,R|I} = \beta_{1,D|R} \quad (14)$$

- In general, with J alternatives, only $J - 1$ binary logits need to be fit (minimal set)

The multinomial logit model: Intuition

VARIABLES	Binary		
	(1) dem_ind	(2) rep_ind	(3) dem_rep
income	-0.00249 (0.00355)	0.0157*** (0.00374)	-0.0184*** (0.00230)
Constant	1.605*** (0.149)	0.659*** (0.162)	0.953*** (0.105)
Observations	844	689	1,231

- Fitting the MNLM by fitting a series of binary logits is not optimal

The multinomial logit model: Intuition

VARIABLES	Binary		
	(1) dem_ind	(2) rep_ind	(3) dem_rep
income	-0.00249 (0.00355)	0.0157*** (0.00374)	-0.0184*** (0.00230)
Constant	1.605*** (0.149)	0.659*** (0.162)	0.953*** (0.105)
Observations	844	689	1,231

- ▶ Fitting the MNLM by fitting a series of binary logits is not optimal
 - ▶ Binary logit is based on a different sample.
 - ▶ It ignores the restrictions that are implicit in the definition of the MNLM

The multinomial logit model: formal statement

- Formally

$$\ln \Omega_{m|b}(X) = \ln \frac{Pr(y = m|X)}{Pr(y = b|X)} = X\beta_{m|b} \text{ for } m = 1, \dots, J \quad (15)$$

- where b is the base outcome (reference category)
- These J equations can be solved to compute the probabilities for each outcome

$$Pr(y = m|X) = \frac{\exp(X\beta_{m|b})}{\sum_{j=1}^J \exp(X\beta_{j|b})} \quad (16)$$

The multinomial logit model: Intuition

VARIABLES	Binary			Multinomial Logit		
	(1) dem_ind	(2) rep_ind	(3) dem_rep	Democrat	Independent	Republican
income	-0.00249 (0.00355)	0.0157*** (0.00374)	-0.0184*** (0.00230)	-0.00272 (0.00372)		0.0152*** (0.00366)
Constant	1.605*** (0.149)	0.659*** (0.162)	0.953*** (0.105)	1.613*** (0.153)		0.678*** (0.160)
Observations	844	689	1,231	1,382	1,382	1,382